

Report for MapReduce2 Java

HAN Duqing & GAO Xin

1.8 Remarkable trees of Paris

- List all jobs.

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar listjob
Unknown program 'listjob' chosen.
Valid program names are:
  districts: A map/reduce program that lists the districts in the trees.csv.
  mosttree: A map/reduce program that displays the district that contains the most trees in the trees.csv.
  oldesttree: A map/reduce program that displays the district where the oldest tree is in the trees.csv.
  sortheight: A map/reduce program that sorts the trees height from smallest to largest in the trees.csv.
  speciecount: A map/reduce program that counts the species in the trees.csv.
  speciemaxheight: A map/reduce program that calculates the height of the tallest tree of each
specie in the trees.csv.
  species: A map/reduce program that lists the species in the trees.csv.
  wordcount: A map/reduce program that counts the words in the input files.
```

1.8.1 Districts containing trees

A map/reduce program that lists the districts in the trees.csv.

	Mapper	Reducer
ClassName	DistrictsMapper	DistrictsReducer
InputKeyFormat	LongWritable	Text
InputValueFormat	Text	NullWritable
OutputKeyFormat	Text	Text
OutputValueFormat	NullWritable	NullWritable

- Run the job "districts".

```
yarn jar q7.jar districts trees.csv q1
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar districts trees.csv q1
20/11/07 12:28:27 INFO client.AHSPProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 12:28:27 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748507917, maxDate=1605353307917, sequenceNumber=4672, masterKeyId=42 on ha-hdfs:efrei
20/11/07 12:28:27 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748507917, maxDate=1605353307917, sequenceNumber=4672, masterKeyId=42)
20/11/07 12:28:28 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2480
20/11/07 12:28:28 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 12:28:28 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 12:28:29 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2480
20/11/07 12:28:29 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748507917, maxDate=1605353307917, sequenceNumber=4672, masterKeyId=42)]
20/11/07 12:28:29 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-152/0/resource-types.xml
20/11/07 12:28:29 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 12:28:29 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2480
20/11/07 12:28:29 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_1603290159664_2480/
20/11/07 12:28:29 INFO mapreduce.Job: Running job: job_1603290159664_2480
20/11/07 12:28:40 INFO mapreduce.Job: Job job_1603290159664_2480 running in uber mode : false
20/11/07 12:28:40 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 12:28:49 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 12:28:58 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 12:28:58 INFO mapreduce.Job: Job job_1603290159664_2480 completed successfully
20/11/07 12:28:58 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=84
    FILE: Number of bytes written=493223
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=44
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=20127
    Total time spent by all reduces in occupied slots (ms)=27272
    Total time spent by all map tasks (ms)=6709
    Total time spent by all reduce tasks (ms)=6818
    Total vcore-milliseconds taken by all map tasks=6709
    Total vcore-milliseconds taken by all reduce tasks=6818
    Total megabyte-milliseconds taken by all map tasks=10305024
    Total megabyte-milliseconds taken by all reduce tasks=13963264
  Map-Reduce Framework
    Map input records=98
    Map output records=97
    Map output bytes=277
    Map output materialized bytes=84
    Input split bytes=97
    Combine input records=97
    Combine output records=17
    Reduce input groups=17
    Reduce shuffle bytes=84
    Reduce input records=17
    Reduce output records=17
    Spilled Records=34
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=197
    CPU time spent (ms)=3330
    Physical memory (bytes) snapshot=1459490816
    Virtual memory (bytes) snapshot=7272185856
    Total committed heap usage (bytes)=1556086784
    Peak Map Physical memory (bytes)=1159405568
    Peak Map Virtual memory (bytes)=3394650112
    Peak Reduce Physical memory (bytes)=300085248
    Peak Reduce Virtual memory (bytes)=3877535744
```

```

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=16778
File Output Format Counters
  Bytes Written=44
[dhan@hadoop-edge01 ~]$

```

- Display the result of job "districts".

```
hdfs dfs -cat q1/part-r-00000
```

```

[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q1/part-r-00000
11
12
13
14
15
16
17
18
19
20
3
4
5
6
7
8
9

```

1.8.2 Show all existing species

A map/reduce program that lists the species in the trees.csv.

	Mapper	Reducer
ClassName	SpeciesMapper	SpeciesReducer
InputKeyFormat	LongWritable	Text
InputValueFormat	Text	NullWritable
OutputKeyFormat	Text	Text
OutputValueFormat	NullWritable	NullWritable

- Run the job "species".

```
yarn jar q7.jar species trees.csv q2
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar species trees.csv q2
20/11/07 12:24:26 INFO client.AHSPProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 12:24:26 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748266975, maxDate=1605353066975, sequenceNumber=4671, masterKeyId=42 on ha-hdfs:efrei
20/11/07 12:24:27 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748266975, maxDate=1605353066975, sequenceNumber=4671, masterKeyId=42)
20/11/07 12:24:27 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2479
20/11/07 12:24:27 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 12:24:27 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 12:24:28 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2479
20/11/07 12:24:28 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748266975, maxDate=1605353066975, sequenceNumber=4671, masterKeyId=42)]
20/11/07 12:24:28 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-152/0/resource-types.xml
20/11/07 12:24:28 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 12:24:28 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2479
20/11/07 12:24:28 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_1603290159664_2479/
20/11/07 12:24:28 INFO mapreduce.Job: Running job: job_1603290159664_2479
20/11/07 12:24:39 INFO mapreduce.Job: Job job_1603290159664_2479 running in uber mode : false
20/11/07 12:24:39 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 12:24:48 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 12:24:58 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 12:24:59 INFO mapreduce.Job: Job job_1603290159664_2479 completed successfully
20/11/07 12:24:59 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=395
    FILE: Number of bytes written=493829
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=317
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=21540
    Total time spent by all reduces in occupied slots (ms)=28380
    Total time spent by all map tasks (ms)=7180
    Total time spent by all reduce tasks (ms)=7095
    Total vcore-milliseconds taken by all map tasks=7180
    Total vcore-milliseconds taken by all reduce tasks=7095
    Total megabyte-milliseconds taken by all map tasks=11028480
    Total megabyte-milliseconds taken by all reduce tasks=14530560
  Map-Reduce Framework
    Map input records=98
    Map output records=97
    Map output bytes=835
    Map output materialized bytes=395
    Input split bytes=97
    Combine input records=97
    Combine output records=36
    Reduce input groups=36
    Reduce shuffle bytes=395
    Reduce input records=36
    Reduce output records=36
    Spilled Records=72
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=201
    CPU time spent (ms)=3550
    Physical memory (bytes) snapshot=1456304128
    Virtual memory (bytes) snapshot=7266349056
    Total committed heap usage (bytes)=1562378240
    Peak Map Physical memory (bytes)=1157955584
    Peak Map Virtual memory (bytes)=3393851392
    Peak Reduce Physical memory (bytes)=298348544
    Peak Reduce Virtual memory (bytes)=3872497664
```

```
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=16778
File Output Format Counters
    Bytes Written=317
[dhan@hadoop-edge01 ~]$
```

- Display the result of job "species".

```
hdfs dfs -cat q2/part-r-00000
```

```
[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q2/part-r-00000
Acer
Aesculus
Ailanthus
Alnus
Araucaria
Broussonetia
Calocedrus
Catalpa
Cedrus
Celtis
Corylus
Davidia
Diospyros
Eucommia
Fagus
Fraxinus
Ginkgo
Gymnocladus
Juglans
Liriodendron
Maclura
Magnolia
Paulownia
Pinus
Platanus
Pterocarya
Quercus
Robinia
Sequoia
Sequoiadendron
Styphnolobium
Taxodium
Taxus
Tilia
Ulmus
Zelkova
[dhan@hadoop-edge01 ~]$
```

1.8.3 Number of trees by species

A map/reduce program that counts the species in the trees.csv.

	Mapper	Reducer
ClassName	SpecieCountMapper	SpecieCountReducer
InputKeyFormat	LongWritable	Text
InputValueFormat	Text	IntWritable
OutputKeyFormat	Text	Text
OutputValueFormat	IntWritable	IntWritable

- Run the job "speciecount".

```
yarn jar q7.jar speciecount trees.csv q3
```



```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar speciecount trees.csv q3
20/11/07 12:30:42 INFO client.AHSPProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 12:30:42 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748642253, maxDate=1605353442253, sequenceNumber=4673, masterKeyId=42 on ha-hdfs:efrei
20/11/07 12:30:42 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748642253, maxDate=1605353442253, sequenceNumber=4673, masterKeyId=42)
20/11/07 12:30:42 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2481
20/11/07 12:30:43 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 12:30:43 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 12:30:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2481
20/11/07 12:30:43 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604748642253, maxDate=1605353442253, sequenceNumber=4673, masterKeyId=42)]
20/11/07 12:30:43 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-152/0/resource-types.xml
20/11/07 12:30:43 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 12:30:44 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2481
20/11/07 12:30:44 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_1603290159664_2481/
20/11/07 12:30:44 INFO mapreduce.Job: Running job: job_1603290159664_2481
20/11/07 12:30:54 INFO mapreduce.Job: Job job_1603290159664_2481 running in uber mode : false
20/11/07 12:30:54 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 12:31:03 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 12:31:13 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 12:31:13 INFO mapreduce.Job: Job job_1603290159664_2481 completed successfully
20/11/07 12:31:13 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=539
    FILE: Number of bytes written=494147
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=390
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=21216
    Total time spent by all reduces in occupied slots (ms)=27632
    Total time spent by all map tasks (ms)=7072
    Total time spent by all reduce tasks (ms)=6908
    Total vcore-milliseconds taken by all map tasks=7072
    Total vcore-milliseconds taken by all reduce tasks=6908
    Total megabyte-milliseconds taken by all map tasks=10862592
    Total megabyte-milliseconds taken by all reduce tasks=14147584
  Map-Reduce Framework
    Map input records=98
    Map output records=97
    Map output bytes=1223
    Map output materialized bytes=539
    Input split bytes=97
    Combine input records=97
    Combine output records=36
    Reduce input groups=36
    Reduce shuffle bytes=539
    Reduce input records=36
    Reduce output records=36
    Spilled Records=72
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=205
    CPU time spent (ms)=3290
    Physical memory (bytes) snapshot=1457258496
    Virtual memory (bytes) snapshot=7273926656
    Total committed heap usage (bytes)=1560805376
    Peak Map Physical memory (bytes)=1156091904
    Peak Map Virtual memory (bytes)=3393040384
    Peak Reduce Physical memory (bytes)=301166592
    Peak Reduce Virtual memory (bytes)=3880886272
```

```

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=16778
File Output Format Counters
  Bytes Written=390
[dhan@hadoop-edge01 ~]$

```

- Display the result of job "speciecount" (key-value: specie-count).

```
hdfs dfs -cat q3/part-r-00000
```

```

[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q3/part-r-00000
Acer      3
Aesculus  3
Ailanthus 1
Alnus     1
Araucaria 1
Broussonetia 1
Calocedrus 1
Catalpa 1
Cedrus    4
Celtis    1
Corylus   3
Davidia   1
Diospyros 4
Eucommia  1
Fagus     8
Fraxinus  1
Ginkgo    5
Gymnocladus 1
Juglans   1
Liriodendron 2
Maclura   1
Magnolia  1
Paulownia 1
Pinus     5
Platanus  19
Pterocarya 3
Quercus   4
Robinia   1
Sequoia   1
Sequoiadendron 5
Styphnolobium 1
Taxodium  3
Taxus     2
Tilia     1
Ulmus     1
Zelkova   4
[dhan@hadoop-edge01 ~]$

```

1.8.4 Maximum height per specie of tree

A map/reduce program that calculates the height of the tallest tree of each specie in the trees.csv.

	Mapper	Reducer
ClassName	SpecieMaxHeightMapper	SpecieMaxHeightReducer
InputKeyFormat	LongWritable	Text
InputValueFormat	Text	FloatWritable
OutputKeyFormat	Text	Text
OutputValueFormat	FloatWritable	FloatWritable

- Run the job "speciemaxheight".

```
yarn jar q7.jar speciemaxheight trees.csv q4
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar speciemaxheight trees.csv q4
20/11/07 12:33:55 INFO client.AHSProxy: Connecting to Application History server at hadoop-master03.efrei.c
nline/163.172.100.24:10200
20/11/07 12:33:55 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLIN
E, renewer=yarn, realUser=, issueDate=1604748835945, maxDate=1605353635945, sequenceNumber=4674, masterKeyId=42 on ha-hdfs:efrei
20/11/07 12:33:55 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service:
ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUse
r=, issueDate=1604748835945, maxDate=1605353635945, sequenceNumber=4674, masterKeyId=42)
20/11/07 12:33:56 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging
/job_1603290159664_2482
20/11/07 12:33:56 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 12:33:57 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 12:33:57 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2482
20/11/07 12:33:57 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service:
: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realU
ser=, issueDate=1604748835945, maxDate=1605353635945, sequenceNumber=4674, masterKeyId=42)]
20/11/07 12:33:57 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-15
2/0/resource-types.xml
20/11/07 12:33:57 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 12:33:57 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2482
20/11/07 12:33:57 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/p
roxy/application_1603290159664_2482/
20/11/07 12:33:57 INFO mapreduce.Job: Running job: job_1603290159664_2482
20/11/07 12:34:08 INFO mapreduce.Job: Job job_1603290159664_2482 running in uber mode : false
20/11/07 12:34:08 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 12:34:17 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 12:34:22 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 12:34:22 INFO mapreduce.Job: Job job_1603290159664_2482 completed successfully
20/11/07 12:34:22 INFO mapreduce.Job: Counters: 53

File System Counters
    FILE: Number of bytes read=539
    FILE: Number of bytes written=494183
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=496
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2

Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=19113
    Total time spent by all reduces in occupied slots (ms)=11792
    Total time spent by all map tasks (ms)=6371
    Total time spent by all reduce tasks (ms)=2948
    Total vcore-milliseconds taken by all map tasks=6371
    Total vcore-milliseconds taken by all reduce tasks=2948
    Total megabyte-milliseconds taken by all map tasks=9785856
    Total megabyte-milliseconds taken by all reduce tasks=6037504

Map-Reduce Framework
    Map input records=98
    Map output records=96
    Map output bytes=1204
    Map output materialized bytes=539
    Input split bytes=97
    Combine input records=96
    Combine output records=36
    Reduce input groups=36
    Reduce shuffle bytes=539
    Reduce input records=36
    Reduce output records=36
    Spilled Records=72
    Shuffled Maps=1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=183
    CPU time spent (ms)=3140
    Physical memory (bytes) snapshot=1452351488
    Virtual memory (bytes) snapshot=7267012608
    Total committed heap usage (bytes)=1570766848
    Peak Map Physical memory (bytes)=1158815744
    Peak Map Virtual memory (bytes)=3393449984
    Peak Reduce Physical memory (bytes)=293535744
    Peak Reduce Virtual memory (bytes)=3873562624
```

```

Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=16778
File Output Format Counters
    Bytes Written=496
[dhan@hadoop-edge01 ~]$

```

- Display the result of job "speciemaxheight" (key-value: specie-maxheight).

```
hdfs dfs -cat q4/part-r-00000
```

```

[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q4/part-r-00000
Acer      16.0
Aesculus  30.0
Ailanthus 35.0
Alnus     16.0
Araucaria 9.0
Broussonetia 12.0
Calocedrus 20.0
Catalpa 15.0
Cedrus    30.0
Celtis    16.0
Corylus   20.0
Davidia   12.0
Diospyros 14.0
Eucommia  12.0
Fagus     30.0
Fraxinus  30.0
Ginkgo    33.0
Gymnocladus 10.0
Juglans   28.0
Liriodendron 35.0
Maclura   13.0
Magnolia  12.0
Paulownia 20.0
Pinus     30.0
Platanus  45.0
Pterocarya 30.0
Quercus   31.0
Robinia   11.0
Sequoia   30.0
Sequoiadendron 35.0
Styphnolobium 10.0
Taxodium  35.0
Taxus     13.0
Tilia     20.0
Ulmus     15.0
Zelkova   30.0
[dhan@hadoop-edge01 ~]$

```

1.8.5 Sort the trees height from smallest to largest

A map/reduce program that sorts the trees height from smallest to largest in the trees.csv.

	Mapper	Reducer
ClassName	SortHeightMapper	SortHeightReducer
InputKeyFormat	LongWritable	FloatWritable
InputValueFormat	Text	Text
OutputKeyFormat	FloatWritable	FloatWritable
OutputValueFormat	Text	Text

- Run the job "sortheight".

```
yarn jar q7.jar sortheight trees.csv q5
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar sortheight trees.csv q5
20/11/07 12:37:02 INFO client.AHSPProxy: Connecting to Application History server at hadoop-master03.efrei.c
nline/163.172.100.24:10200
20/11/07 12:37:02 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLIN
E, renewer=yarn, realUser=, issueDate=1604749022353, maxDate=1605353822353, sequenceNumber=4675, masterKeyId
d=42 on ha-hdfs:efrei
20/11/07 12:37:02 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service:
ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUse
r=, issueDate=1604749022353, maxDate=1605353822353, sequenceNumber=4675, masterKeyId=42)
20/11/07 12:37:02 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/
g/job_1603290159664_2483
20/11/07 12:37:03 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 12:37:03 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 12:37:03 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2483
20/11/07 12:37:03 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service
: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realU
ser=, issueDate=1604749022353, maxDate=1605353822353, sequenceNumber=4675, masterKeyId=42)]
20/11/07 12:37:03 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-15
2/0/resource-types.xml
20/11/07 12:37:03 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 12:37:04 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2483
20/11/07 12:37:04 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/p
roxy/application_1603290159664_2483/
20/11/07 12:37:04 INFO mapreduce.Job: Running job: job_1603290159664_2483
20/11/07 12:37:14 INFO mapreduce.Job: Job job_1603290159664_2483 running in uber mode : false
20/11/07 12:37:14 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 12:37:23 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 12:37:31 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 12:37:32 INFO mapreduce.Job: Job job_1603290159664_2483 completed successfully
20/11/07 12:37:32 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=861
    FILE: Number of bytes written=494787
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=755
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=20157
    Total time spent by all reduces in occupied slots (ms)=26596
    Total time spent by all map tasks (ms)=6719
    Total time spent by all reduce tasks (ms)=6649
    Total vcore-milliseconds taken by all map tasks=6719
    Total vcore-milliseconds taken by all reduce tasks=6649
    Total megabyte-milliseconds taken by all map tasks=10320384
    Total megabyte-milliseconds taken by all reduce tasks=13617152
  Map-Reduce Framework
    Map input records=98
    Map output records=96
    Map output bytes=663
    Map output materialized bytes=861
    Input split bytes=97
    Combine input records=96
    Combine output records=96
    Reduce input groups=28
    Reduce shuffle bytes=861
    Reduce input records=96
    Reduce output records=96
    Spilled Records=192
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=197
    CPU time spent (ms)=3220
    Physical memory (bytes) snapshot=1457455104
    Virtual memory (bytes) snapshot=7271501824
    Total committed heap usage (bytes)=1552416768
    Peak Map Physical memory (bytes)=1156358144
    Peak Map Virtual memory (bytes)=3394916352
    Peak Reduce Physical memory (bytes)=301096960
    Peak Reduce Virtual memory (bytes)=3876585472
```

```
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=16778
File Output Format Counters
    Bytes Written=755
[dhan@hadoop-edge01 ~]$
```

- Display the result of job "sortheight" (key-value: height-objectid).

```
hdfs dfs -cat q5/part-r-00000
```



```
[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q5/part-r-00000
```

```
2.0      3
5.0      89
6.0      62
9.0      39
10.0     44
10.0     32
10.0     95
10.0     61
10.0     63
11.0      4
12.0     93
12.0     66
12.0     50
12.0      7
12.0     48
12.0     58
12.0     33
12.0     71
13.0     36
13.0      6
14.0     68
14.0     96
14.0     94
15.0     91
15.0      5
15.0     70
15.0      2
15.0     98
16.0     28
16.0     78
16.0     75
16.0     16
18.0     64
18.0     83
18.0     23
18.0     60
20.0     34
20.0     51
20.0     43
20.0     15
20.0     11
20.0      1
20.0      8
20.0     20
20.0     35
20.0     12
20.0     87
20.0     13
22.0     10
22.0     47
22.0     86
22.0     14
22.0     88
23.0     18
25.0     24
25.0     31
25.0     92
25.0     49
```

```
25.0    97
25.0    84
26.0    73
27.0    65
27.0    42
28.0    85
30.0    76
30.0    19
30.0    72
30.0    54
30.0    29
30.0    27
30.0    25
30.0    41
30.0    77
30.0    55
30.0    69
30.0    38
30.0    59
30.0    52
30.0    37
30.0    22
30.0    30
31.0    80
31.0     9
32.0    82
33.0    46
34.0    45
35.0    56
35.0    81
35.0    17
35.0    53
35.0    57
40.0    26
40.0    74
40.0    40
42.0    90
45.0    21
[dhan@hadoop-edge01 ~]$
```

1.8.6 District containing the oldest tree

A map/reduce program that displays the district where the oldest tree is in the trees.csv.

We define a Writable class called MyWritable which is the pair of (plantyear , district).

	Mapper	Reducer
ClassName	OldestTreeMapper	OldestTreeReducer
InputKeyFormat	LongWritable	Text
InputValueFormat	Text	MyWritable
OutputKeyFormat	Text	Text
OutputValueFormat	MyWritable	NullWritable

- Run the job "oldesttree" and display the result.

```
yarn jar q7.jar oldesttree trees.csv q6
hdfs dfs -cat q6/part-r-00000
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q6.jar oldesttree trees.csv q6
20/11/07 14:29:57 INFO client.AHSProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 14:29:57 INFO hdfs.DFSCClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, i
ssueDate=1604755797815, maxDate=1605360597815, sequenceNumber=4709, masterKeyId=42 on ha-hdfs:efrei
20/11/07 14:29:57 INFO security.TokenCache: Got dt for hdfs://efrei: Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token
for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755797815, maxDate=1605360597815, sequen
ceNumber=4709, masterKeyId=42)
20/11/07 14:29:57 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2515
20/11/07 14:29:58 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 14:29:58 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 14:29:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2515
20/11/07 14:29:58 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token
n for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755797815, maxDate=1605360597815, sequ
enceNumber=4709, masterKeyId=42)]
20/11/07 14:29:59 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-152/0/resource-types.xml
20/11/07 14:29:59 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 14:29:59 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2515
20/11/07 14:29:59 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_160329015966
4_2515/
20/11/07 14:29:59 INFO mapreduce.Job: Running job: job_1603290159664_2515
20/11/07 14:30:09 INFO mapreduce.Job: Job job_1603290159664_2515 running in uber mode : false
20/11/07 14:30:09 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 14:30:19 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 14:30:29 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 14:30:30 INFO mapreduce.Job: Job job_1603290159664_2515 completed successfully
20/11/07 14:30:30 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=1072
    FILE: Number of bytes written=495575
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=2
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=21747
    Total time spent by all reduces in occupied slots (ms)=28528
    Total time spent by all map tasks (ms)=7249
    Total time spent by all reduce tasks (ms)=7132
    Total vcore-milliseconds taken by all map tasks=7249
    Total vcore-milliseconds taken by all reduce tasks=7132
    Total megabyte-milliseconds taken by all map tasks=11134464
    Total megabyte-milliseconds taken by all reduce tasks=14606336
  Map-Reduce Framework
    Map input records=98
    Map output records=77
    Map output bytes=912
    Map output materialized bytes=1072
    Input split bytes=97
    Combine input records=0
    Combine output records=0
    Reduce input groups=1
    Reduce shuffle bytes=1072
    Reduce input records=77
    Reduce output records=1
    Spilled Records=154
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=196
    CPU time spent (ms)=3450
    Physical memory (bytes) snapshot=1453199360
    Virtual memory (bytes) snapshot=7272005632
    Total committed heap usage (bytes)=1551368192
    Peak Map Physical memory (bytes)=1159045120
    Peak Map Virtual memory (bytes)=3395592192
    Peak Reduce Physical memory (bytes)=294154240
    Peak Reduce Virtual memory (bytes)=3876413440
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=16778
  File Output Format Counters
    Bytes Written=2
[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q6/part-r-00000
5
[dhan@hadoop-edge01 ~]$
```

- Verify the result of job "oldesttree".

```
df[df['ANNEE PLANTATION']==min(df['ANNEE PLANTATION'])]
```

	GEOPOINT	ARRONDISSEMENT	GENRE	ESPECE	FAMILLE	ANNEE PLANTATION	HAUTEUR	CIRCONFERENCE	ADRESSE	NOM COMMUN	VARIETE	OBJEC
26	(48.8520958913, 2.34740754195)	5	Robinia	pseudoacacia	Fabaceae	1601.0	11.0	385.0	Quai de Montebello, rue Lagrange, rue Saint-Ju...	Robinier faux-acacia	NaN	

We can see that District 5 have the oldest tree is correct.

1.8.7 District containing the most trees

A map/reduce program that displays the district that contains the most trees in the trees.csv.

We define a Writable class called MyWritable which is the pair of (district , count).

We also create 2 MapReduce jobs to handle this question.

	Mapper1	Reducer1	Mapper2	Reducer2
ClassName	MostTreeMapper1	MostTreeReducer1	MostTreeMapper2	MostTreeReducer2
InputKeyFormat	LongWritable	Text	Oject	NullWritable
InputValueFormat	Text	IntWritable	Text	MyWritable
OutputKeyFormat	Text	Text	NullWritable	Text
OutputValueFormat	IntWritable	IntWritable	MyWritable	NullWritable

- Run the job "mosttree" and display the result.

```
yarn jar q7.jar mosttree trees.csv q7
hdfs dfs -cat q7/part-r-00000
```

```
[dhan@hadoop-edge01 ~]$ yarn jar q7.jar mosttree trees.csv q7
20/11/07 14:20:50 INFO client.AHSProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 14:20:50 INFO hdfs.DFSClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755250568, maxDate=1605360050568, sequenceNumber=4705, masterKeyId=42 on ha-hdfs:efrei
20/11/07 14:20:50 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755250568, maxDate=1605360050568, sequenceNumber=4705, masterKeyId=42)
20/11/07 14:20:50 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2511
20/11/07 14:20:51 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 14:20:51 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 14:20:51 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2511
20/11/07 14:20:51 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755250568, maxDate=1605360050568, sequenceNumber=4705, masterKeyId=42)]
20/11/07 14:20:51 INFO conf.Configuration: found resource resource-types.xml at file:/etc/hadoop/3.1.5.0-152/0/resource-types.xml
20/11/07 14:20:51 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 14:20:52 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2511
20/11/07 14:20:52 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_1603290159664_2511/
20/11/07 14:20:52 INFO mapreduce.Job: Running job: job_1603290159664_2511
20/11/07 14:21:02 INFO mapreduce.Job: Job job_1603290159664_2511 running in uber mode : false
20/11/07 14:21:02 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 14:21:11 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 14:21:20 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 14:21:21 INFO mapreduce.Job: Job job_1603290159664_2511 completed successfully
20/11/07 14:21:22 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=865
    FILE: Number of bytes written=495181
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=16875
    HDFS: Number of bytes written=80
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=20148
    Total time spent by all reduces in occupied slots (ms)=28080
    Total time spent by all map tasks (ms)=6716
    Total time spent by all reduce tasks (ms)=7020
    Total vcore-milliseconds taken by all map tasks=6716
    Total vcore-milliseconds taken by all reduce tasks=7020
    Total megabyte-milliseconds taken by all map tasks=10315776
    Total megabyte-milliseconds taken by all reduce tasks=14376960
  Map-Reduce Framework
    Map input records=98
    Map output records=97
    Map output bytes=665
    Map output materialized bytes=865
    Input split bytes=97
    Combine input records=0
    Combine output records=0
    Reduce input groups=17
    Reduce shuffle bytes=865
    Reduce input records=97
    Reduce output records=17
    Spilled Records=194
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=185
    CPU time spent (ms)=3560
    Physical memory (bytes) snapshot=1453039616
    Virtual memory (bytes) snapshot=7273865216
    Total committed heap usage (bytes)=1571815424
    Peak Map Physical memory (bytes)=1158545408
    Peak Map Virtual memory (bytes)=3397316608
    Peak Reduce Physical memory (bytes)=294494208
    Peak Reduce Virtual memory (bytes)=3876548608
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=16778
  File Output Format Counters
    Bytes Written=80
20/11/07 14:21:22 INFO client.AHSProxy: Connecting to Application History server at hadoop-master03.efrei.online/163.172.100.24:10200
20/11/07 14:21:22 INFO hdfs.DFSClient: Created token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755282102, maxDate=1605360082102, sequenceNumber=4706, masterKeyId=42 on ha-hdfs:efrei
20/11/07 14:21:22 INFO security.TokenCache: Got dt for hdfs://efrei; Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident: (token for dhan: HDFS_DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755282102, maxDate=1605360082102, sequenceNumber=4706, masterKeyId=42)
20/11/07 14:21:22 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/11/07 14:21:22 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /user/dhan/.staging/job_1603290159664_2512
```

```

20/11/07 14:21:22 INFO input.FileInputFormat: Total input files to process : 1
20/11/07 14:21:22 INFO mapreduce.JobSubmitter: number of splits:1
20/11/07 14:21:22 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1603290159664_2512
20/11/07 14:21:22 INFO mapreduce.JobSubmitter: Executing with tokens: [Kind: HDFS_DELEGATION_TOKEN, Service: ha-hdfs:efrei, Ident:
(token for dhan: HDFS DELEGATION_TOKEN owner=dhan@EFREI.ONLINE, renewer=yarn, realUser=, issueDate=1604755282102, maxDate=160536008
2102, sequenceNumber=4706, masterKeyId=42)]
20/11/07 14:21:22 INFO impl.TimelineClientImpl: Timeline service address: hadoop-master03.efrei.online:8190
20/11/07 14:21:23 INFO impl.YarnClientImpl: Submitted application application_1603290159664_2512
20/11/07 14:21:23 INFO mapreduce.Job: The url to track the job: https://hadoop-master01.efrei.online:8090/proxy/application_1603290
159664_2512/
20/11/07 14:21:23 INFO mapreduce.Job: Running job: job_1603290159664_2512
20/11/07 14:21:33 INFO mapreduce.Job: Job job_1603290159664_2512 running in uber mode : false
20/11/07 14:21:33 INFO mapreduce.Job: map 0% reduce 0%
20/11/07 14:21:42 INFO mapreduce.Job: map 100% reduce 0%
20/11/07 14:21:52 INFO mapreduce.Job: map 100% reduce 100%
20/11/07 14:21:52 INFO mapreduce.Job: Job job_1603290159664_2512 completed successfully
20/11/07 14:21:52 INFO mapreduce.Job: Counters: 53
  File System Counters
    FILE: Number of bytes read=169
    FILE: Number of bytes written=493427
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=192
    HDFS: Number of bytes written=3
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=21090
    Total time spent by all reduces in occupied slots (ms)=26812
    Total time spent by all map tasks (ms)=7030
    Total time spent by all reduce tasks (ms)=6703
    Total vcore-milliseconds taken by all map tasks=7030
    Total vcore-milliseconds taken by all reduce tasks=6703
    Total megabyte-milliseconds taken by all map tasks=10798080
    Total megabyte-milliseconds taken by all reduce tasks=13727744
  Map-Reduce Framework
    Map input records=17
    Map output records=17
    Map output bytes=129
    Map output materialized bytes=169
    Input split bytes=112
    Combine input records=0
    Combine output records=0
    Reduce input groups=1
    Reduce shuffle bytes=169
    Reduce input records=17
    Reduce output records=1
    Spilled Records=34
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=171
    CPU time spent (ms)=3140
    Physical memory (bytes) snapshot=1452314624
    Virtual memory (bytes) snapshot=7268995072
    Total committed heap usage (bytes)=1556611072
    Peak Map Physical memory (bytes)=1163288576
    Peak Map Virtual memory (bytes)=3393609728
    Peak Reduce Physical memory (bytes)=289026048
    Peak Reduce Virtual memory (bytes)=3875385344
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=80
  File Output Format Counters
    Bytes Written=3
[dhan@hadoop-edge01 ~]$ hdfs dfs -cat q7/part-r-00000
16
[dhan@hadoop-edge01 ~]$

```

- Verify the result of job "oldesttree".


```
df.groupby('ARRONDISSEMENT').count()['OBJECTID'].sort_values(ascending=False)
```

ARRONDISSEMENT

16 36

12 29

19 6

8 5

14 3

7 3

20 3

13 2

5 2

15 1

11 1

9 1

17 1

18 1

6 1

4 1

3 1

Name: OBJECTID, dtype: int64

We can see that District 16 have most trees is correct.